

nference-propensity-score-matching

November 10, 2025

1 Step 03: Causal Inference & Uplift Modeling

Project: AI Experimentation Platform for Predictive Insights

Notebook: 03_Causal_Inference_and_Uplift_Modeling

Author: Jagadeesh N

Date: 10-Nov-2025

Environment: Python (DoWhy · EconML · Scikit-learn · Matplotlib · Seaborn)

1.0.1 Objective

This notebook estimates the **true causal impact** of being in the *treatment group (B)* on customer conversion and revenue.

We will move beyond correlation and apply **causal inference** and **uplift modeling** to measure how much the experimental treatment *caused* behavior change.

1.0.2 Background

A/B testing shows performance differences between two groups, but not *why* those differences occur. Causal inference methods (like **DoWhy** and **EconML**) help estimate **treatment effects** while controlling for confounders (e.g., spending, number of transactions, high-value behavior).

1.0.3 Notebook Outline

1. Load simulation dataset & baseline model
2. Estimate Average Treatment Effect (ATE) using DoWhy
3. Conduct causal graph-based validation
4. Apply EconML's Double ML method for uplift estimation
5. Visualize treatment effect distribution
6. Interpret business insights

```
[1]: #importing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split

sns.set(style="whitegrid")
np.random.seed(42)

# Load simulated dataset
data = pd.read_csv("D:\\Project for job\\ab_test_simulated.csv")
data.head()
```

	Customer ID	Revenue	NumTransactions	TotalUnits	Converted	Group
0	12346.0	93843.3166	12	74285	0	A
1	12347.0	5955.0513	8	2967	1	B
2	12348.0	2221.3400	5	2714	0	A
3	12349.0	5358.7149	4	1624	0	B
4	12350.0	367.8400	1	197	0	A

```
[2]: # Treatment: being in group B (1)
data['Treatment'] = (data['Group'] == 'B').astype(int)

# Outcome: whether the customer converted
outcome = 'Converted'

# Confounders (features that influence both treatment and outcome)
confounders = [
    'TotalUnits', 'NumTransactions', 'Revenue',
    'RevenuePerUnit', 'RevenuePerTransaction', 'HighValueCustomer'
]

print("Treatment:", "Treatment")
print("Outcome:", outcome)
print("Confounders:", confounders)
```

Treatment: Treatment
 Outcome: Converted
 Confounders: ['TotalUnits', 'NumTransactions', 'Revenue', 'RevenuePerUnit',
 'RevenuePerTransaction', 'HighValueCustomer']

```
[5]: print(data.columns.tolist())
```

['Customer ID', 'Revenue', 'NumTransactions', 'TotalUnits', 'Converted',
 'Group', 'Treatment']

```
[8]: # Create binary treatment variable
data['treatment'] = (data['Group'] == 'B').astype(int)

[9]: data[['Group', 'treatment']].head()

[9]:   Group  treatment
0      A          0
1      B          1
2      A          0
3      B          1
4      A          0

[10]: from dowhy import CausalModel

treatment = 'treatment'                      # just created above
outcome = 'Converted'                        # your binary outcome
confounders = [col for col in data.columns if col not in ['Group', 'Converted', ↴
    'Revenue', 'treatment']]

[12]: model = CausalModel(
        data=data,
        treatment=treatment,
        outcome=outcome,
        common_causes=confounders
    )

    identified_estimand = model.identify_effect()
    estimate = model.estimate_effect(
        identified_estimand,
        method_name="backdoor.linear_regression"
    )

    print(" Estimated Average Treatment Effect (ATE):", estimate.value)

    refutation = model.refute_estimate(
        identified_estimand,
        estimate,
        method_name="placebo_treatment_refuter"
    )
    print(refutation)

Estimated Average Treatment Effect (ATE): 0.009893893320552238
Refute: Use a Placebo Treatment
Estimated effect:0.009893893320552238
New effect:-0.0008547437230948847
p value:0.88
```

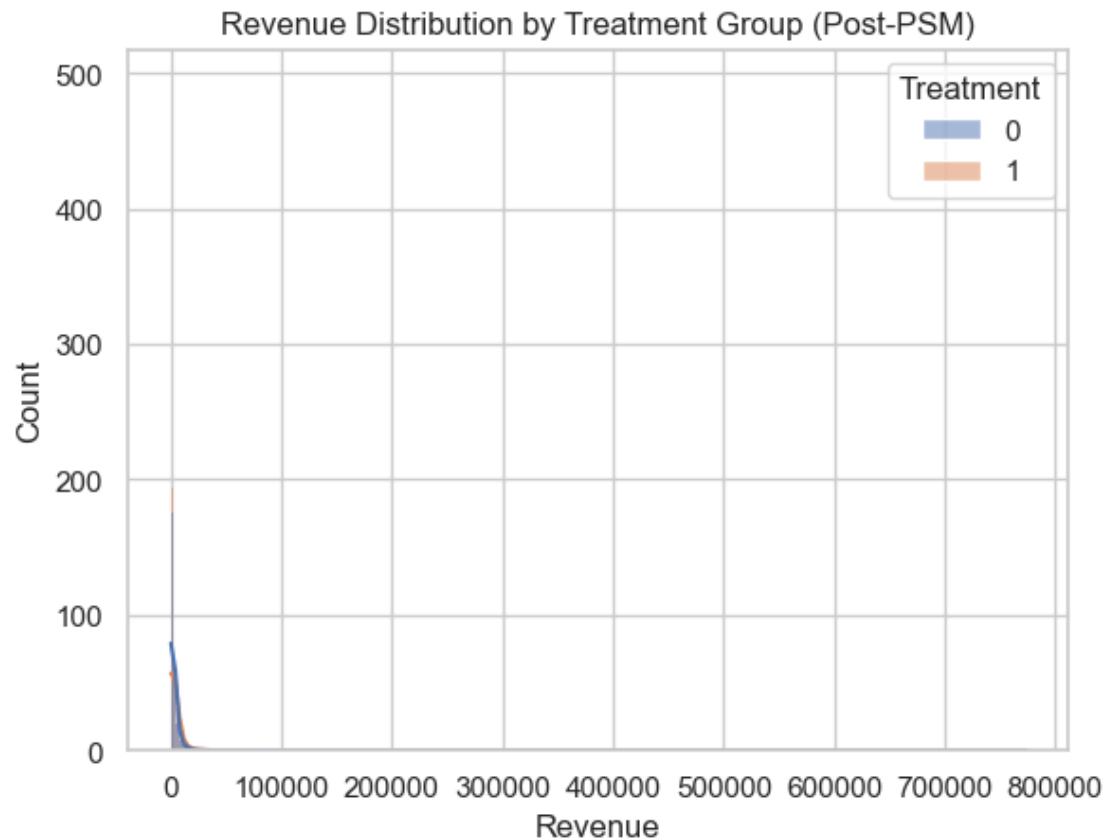
```
[15]: from sklearn.linear_model import LogisticRegression
estimate = model.estimate_effect(
    identified_estimand,
    method_name="backdoor.propensity_score_matching",
    method_params={"propensity_score_model": LogisticRegression(max_iter=1000)}
)
```

```
[18]: #Refute the estimate
refute = model.refute_estimate(
    identified_estimand,
    estimate,
    method_name="placebo_treatment_refuter"
)
print(refute)
```

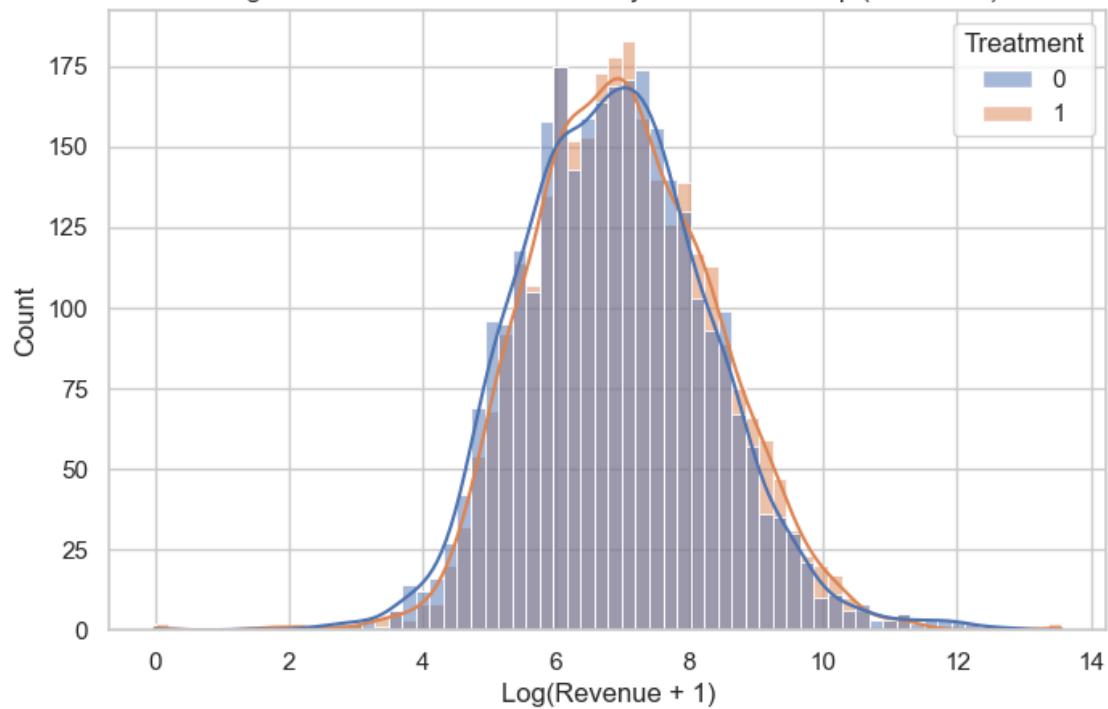
Refute: Use a Placebo Treatment
Estimated effect: 0.009352150994728798
New effect: 0.00021765005951368806
p value: 0.8999999999999999

```
[20]: #visualizing
import matplotlib.pyplot as plt
sns.histplot(data=data, x='Revenue', hue='Treatment', kde=True)
plt.title("Revenue Distribution by Treatment Group (Post-PSM)")
plt.show()
plt.figure(figsize=(8,5))
sns.histplot(
    data=data,
    x=np.log1p(data['Revenue']), # log(Revenue + 1)
    hue='Treatment',
    kde=True
)
plt.title("Log-Scaled Revenue Distribution by Treatment Group (Post-PSM)")
plt.xlabel("Log(Revenue + 1)")
plt.ylabel("Count")
plt.show()
plt.figure(figsize=(8,5))
sns.histplot(
    data=data[data['Revenue'] < data['Revenue'].quantile(0.99)],
    x='Revenue',
    hue='Treatment',
    kde=True
)
plt.title("Revenue Distribution by Treatment Group (Trimmed Top 1%)")
plt.show()
plt.figure(figsize=(6,4))
```

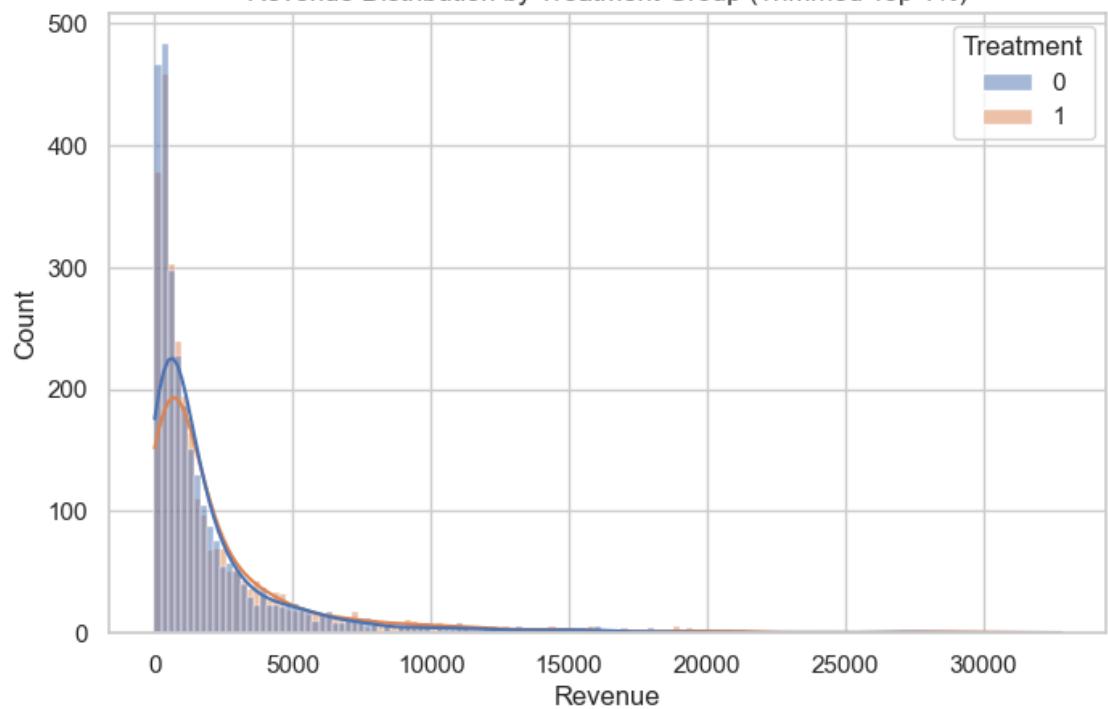
```
sns.boxplot(data=data, x='Treatment', y='Revenue')
plt.title("Revenue Comparison by Treatment (Post-PSM)")
plt.xticks([0, 1], ['Control', 'Treatment'])
plt.show()
```

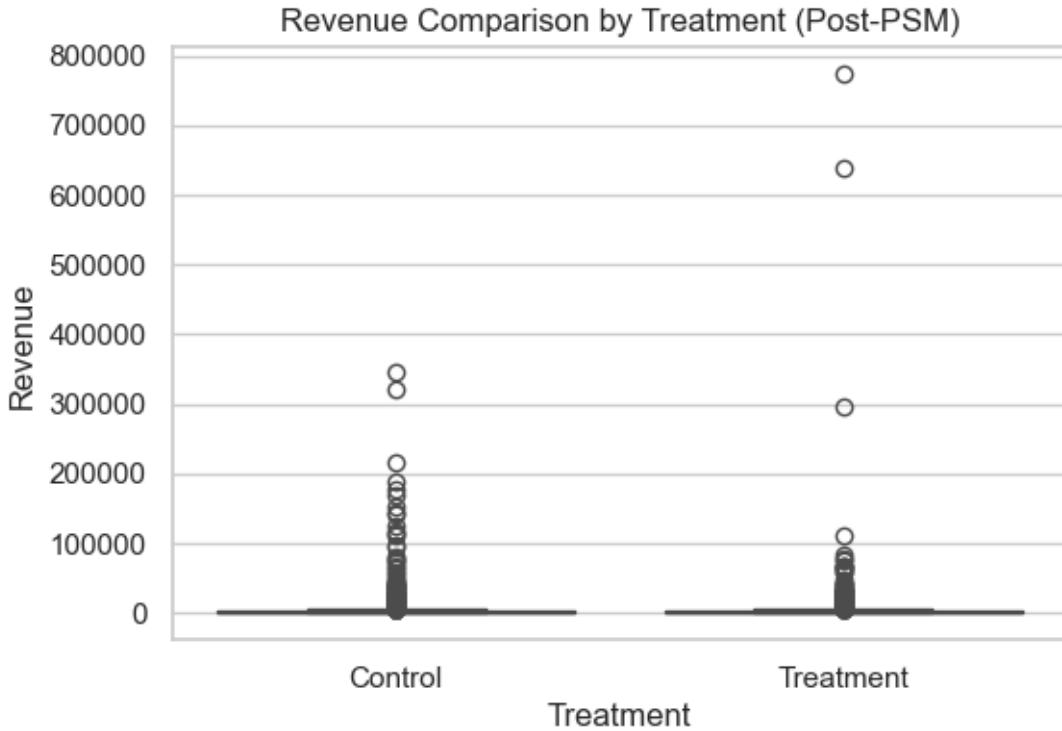


Log-Scaled Revenue Distribution by Treatment Group (Post-PSM)



Revenue Distribution by Treatment Group (Trimmed Top 1%)





```
[23]: original_data = pd.read_csv("D:\\Project for job\\ab_test_simulated.csv")
# AFTER PSM dataset (your current matched data after causal modeling)
matched_data = data.copy()    # 'data' is the one used after PSM in your model
```

```
[26]: print(data.columns.tolist())
```

```
['Customer ID', 'Revenue', 'NumTransactions', 'TotalUnits', 'Converted',
'Group', 'Treatment', 'treatment', 'propensity_score']
```

```
[28]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pandas as pd

# Load original data (before matching)
original_data = pd.read_csv("D:\\Project for job\\ab_test_simulated.csv")

# matched_data is your current 'data' (after PSM)
matched_data = data.copy()

plt.figure(figsize=(14, 5))
fig, axes = plt.subplots(1, 2, figsize=(14, 5))
```

```

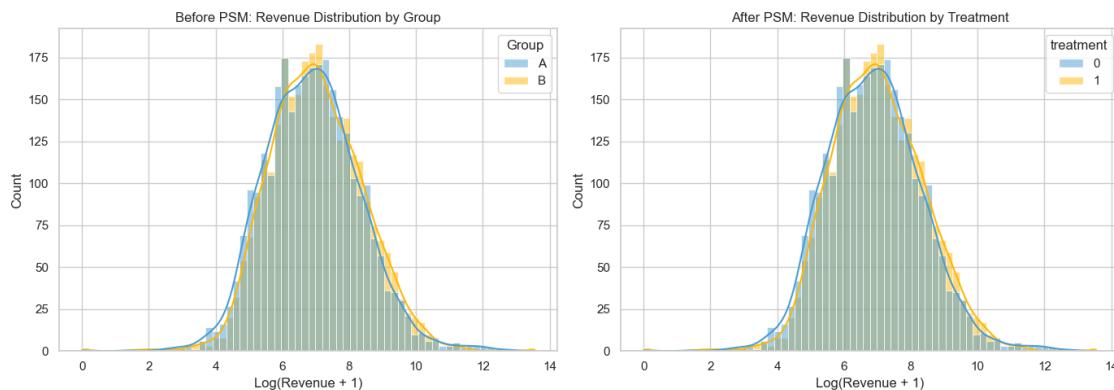
# BEFORE PSM
sns.histplot(
    data=original_data,
    x=np.log1p(original_data['Revenue']),
    hue='Group',
    kde=True,
    ax=axes[0],
    palette=['#4B9CD3', '#FFB703']
)
axes[0].set_title("Before PSM: Revenue Distribution by Group")
axes[0].set_xlabel("Log(Revenue + 1)")
axes[0].set_ylabel("Count")

# AFTER PSM
sns.histplot(
    data=matched_data,
    x=np.log1p(matched_data['Revenue']),
    hue='treatment',
    kde=True,
    ax=axes[1],
    palette=['#4B9CD3', '#FFB703']
)
axes[1].set_title("After PSM: Revenue Distribution by Treatment")
axes[1].set_xlabel("Log(Revenue + 1)")
axes[1].set_ylabel("Count")

plt.tight_layout()
plt.show()

```

<Figure size 1400x500 with 0 Axes>



[31]: `#Save the matched dataset for future steps
matched_data.to_csv("D:\\Project for job\\ab_test_matched.csv", index=False)`

```
print("Matched dataset saved successfully at: D:/Project for job/data/processed/  
ab_test_matched.csv")
```

Matched dataset saved successfully at: D:/Project for
job/data/processed/ab_test_matched.csv